

National Institute of Technology, Srinagar
Department of Information Technology, NIT Srinagar, Hazratbal-190006 Srinagar
Jammu and Kashmir.

DEPARTMENT OF INFORMATION TECHNOLOGY



A REPORT
ON
”Detection, Classification and Semantic
segmentation of apples”

SUBMITTED BY

Mr. TAJAMUL ASHRAF (2018BITE031)
Mr. MOHAMMAD HASEEB (2018BITE024)
Mr. NAIYER ABBAS (2018BITE056)

UNDER THE GUIDANCE OF

PROF. PABITRA MITRA
(IIT KHARAGPUR)

(Internship Period: October 2020 - March 2021)
(Academic Year: 2020-2021)

ACKNOWLEDGEMENT

It gives us great pleasure and satisfaction in presenting this project report on “Detection, Classification and Semantic segmentation of apples”.

We would like to express our deep sense of gratitude towards **Prof. Pabitra Mitra** for providing us this great opportunity to work under his guidance and steering us at every step during this research internship.

*We have furthermore to thank **Mr. Nadeem Yousuf** for encouraging us to go ahead and for his continuous guidance.*

We would like to thank all those, who have directly or indirectly helped us for the completion of the work during this project.

TAJAMUL ASHRAF
MOHAMMAD HASEEB
NAIYER ABBAS

Contents

0.1	Introduction	1
0.2	Problem Statement	1
0.3	Detection and classification of Apples	2
0.3.1	Step1 — Getting Training Data	2
0.3.2	Step2 — Training our model	2
0.3.3	Step 3 — Predicting using our trained Model	4
0.4	Semantic Segmentation	6
0.4.1	U-Net	7
0.4.2	Training	7
0.4.3	Inference	8
0.4.4	Our Results on training set	8

0.1 Introduction

It was a challenging problem that involves building upon methods for apple recognition (e.g. where are they), object localization (e.g. what are their extent), and object classification (e.g. rotten or fresh in our case). After doing some research, we found that in recent years, deep learning techniques are achieving state-of-the-art results for object detection, such as on standard benchmark datasets and in computer vision competitions. Notable is the “You Only Look Once,” or YOLO, family of Convolutional Neural Networks that achieve near state-of-the-art results with a single end-to-end model that can perform object detection in real-time.

We would like to discuss about one specific task in Computer Vision called as Semantic Segmentation. Even though researchers have come up with numerous ways to solve this problem, we worked on a particular architecture namely UNET, which use a Fully Convolutional Network Model for the task.

In this project, we discovered how to develop a model for object detection, classification, counting and semantic segmentation on the given dataset.

<https://arxiv.org/abs/1909.06441>

After completing this project, we are familiar with the following:

- YOLO-based Convolutional Neural Network family of models for object detection and the most recent variation called YOLOv3.
- The best-of-breed open source library implementation of the YOLOv3 for the Keras deep learning library.
- How to use a pre-trained YOLOv3 to perform object localization, detection, classification(rotten / fresh) and counting on new images.
- The goal of semantic image segmentation is to label each pixel of an image with a corresponding class of what is being represented. Because we’re predicting for every pixel in the image, this task is commonly referred to as dense prediction.

0.2 Problem Statement

To use computer vision to aid agricultural robotics.

1. Identification of fruits in a tree
2. Read and understand a fruit detection algorithm using semantic segmentation.
3. Devise an algorithm for counting fruits.
4. Implement the algorithm and run on above data set.
5. Devise an algorithm to identify fresh and rotten fruit..

0.3 Detection and classification of Apples

We used Artificial Intelligence through the use of computer vision and deep learning to automate apple detection, sorting and identification of rotten ones.

0.3.1 Step1 — Getting Training Data

We provided sufficient sample images (dataset) of the apples we need the model to detect and identify. We amended the YOLOV3 algorithm for detecting and counting apples as well as identify rotten apples in images. We used the following dataset:

https://github.com/OlafenwaMoses/AppleDetection/releases/download/v1/apple_detection_dataset.zip

It is divided into:

1. 563 images for training the AI model
2. 150 images for testing the trained AI model

0.3.2 Step2 — Training our model

To generate our model, we trained a YOLOv3 model using ImageAI.

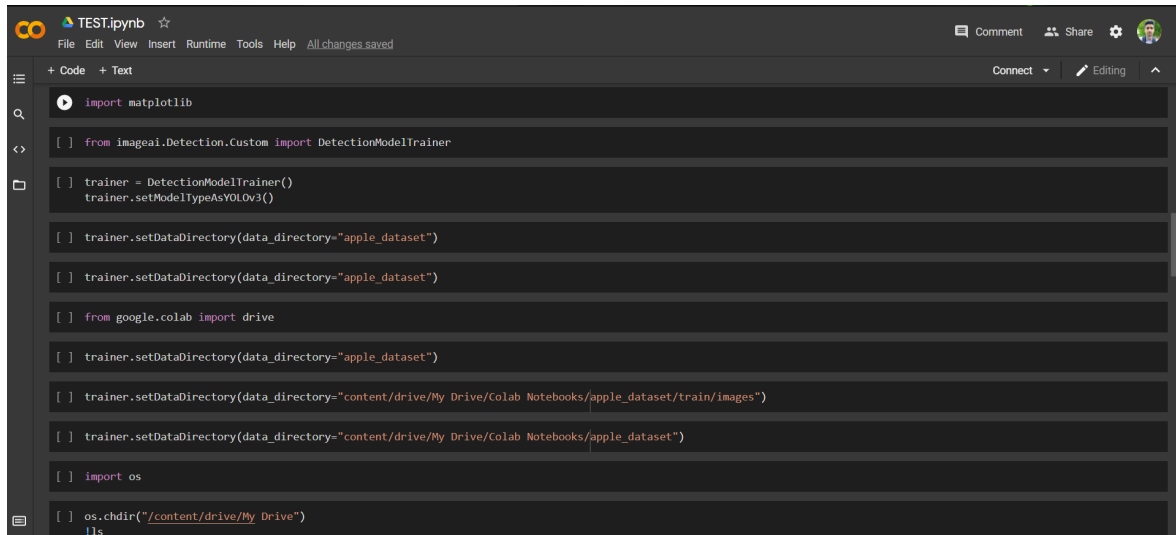
YOLO V3 is designed to be a multi-scaled detector, we also need features from multiple scales. Therefore, features from last three residual blocks are all used in the later detection. We are assuming the input is 416x416, so three scale vectors would be 52x52, 26x26, and 13x13.

Since the convolution outputs a square matrix of feature values (like 13x13, 26x26, and 52x52 in YOLO), we define this matrix as a “grid” and assign anchor boxes to each cell of the grid. In other words, anchor boxes anchor to the grid cells, and they share the same centroid.

With the final detection output, we can calculate the loss against the ground truth labels now. The loss function consists of four parts (or five, if you split noobj and obj): centroid (xy) loss, width and height (wh) loss, objectness (obj and noobj) loss and classification loss.

```
Loss = Lambda_Coord * Sum(Mean_Square_Error((tx, ty),
(tx', ty') * obj_mask)
+ Lambda_Coord * Sum(Mean_Square_Error((tw, th), (tw',
th') * obj_mask)
+ Sum(Binary_Cross_Entropy(obj, obj') * obj_mask) +
Lambda_Noobj * Sum(Binary_Cross_Entropy(obj, obj') * (1
-obj_mask) * ignore_mask)
+ Sum(Binary_Cross_Entropy(class, class'))
```

The below lines of code is to initiate the training on our apple dataset. Now we import the “DetectionModelTrainer” class from ImageAI. After that we created an instance of the class and set our model type to YOLOv3 ,we set the path to our apple dataset and finally we specified the parameters.



```

import matplotlib

[ ] from imageai.Detection.Custom import DetectionModelTrainer

[ ] trainer = DetectionModelTrainer()
    trainer.setModelTypeAsYOLOv3()

[ ] trainer.setDataDirectory(data_directory="apple_dataset")

[ ] trainer.setDataDirectory(data_directory="apple_dataset")

[ ] from google.colab import drive

[ ] trainer.setDataDirectory(data_directory="apple_dataset")

[ ] trainer.setDataDirectory(data_directory="content/drive/My Drive/Colab Notebooks/apple_dataset/train/images")

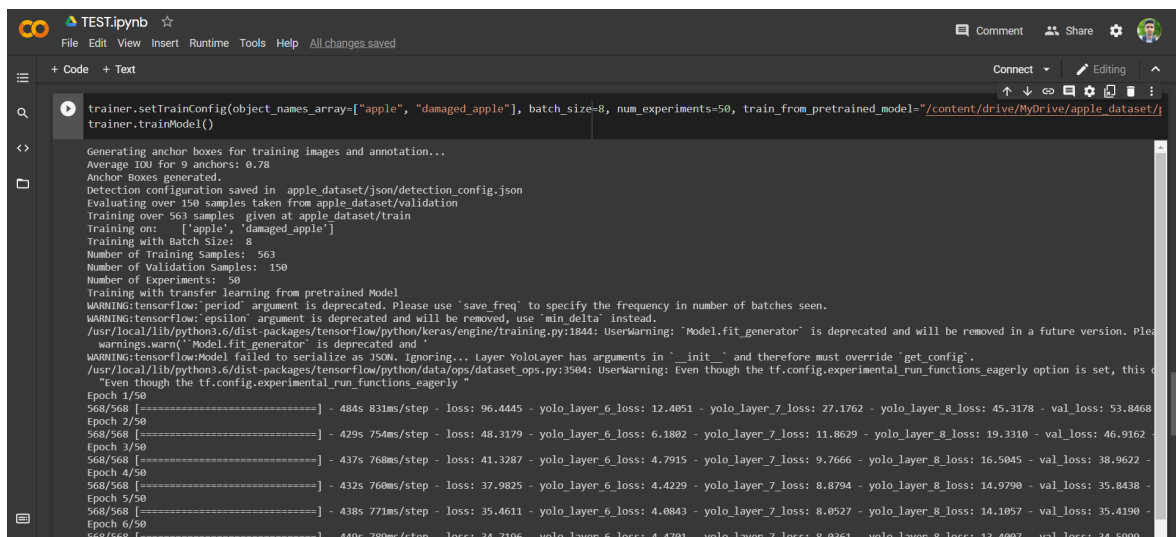
[ ] trainer.setDataDirectory(data_directory="content/drive/My Drive/Colab Notebooks/apple_dataset")

[ ] import os

[ ] os.chdir("/content/drive/My Drive")
11s

```

ImageAI will create apple dataset/models folder which is where all generated models will be saved.



```

trainer.setTrainConfig(object_names_array=["apple", "damaged_apple"], batch_size=8, num_experiments=50, train_from_pretrained_model="/content/drive/MyDrive/apple_dataset/j
trainer.trainModel()

Generating anchor boxes for training images and annotation...
Average IOU for 9 anchors: 0.78
Anchor Boxes generated.
Detection configuration saved in apple_dataset/json/detection_config.json
Evaluating over 150 samples taken from apple_dataset/validation
Training over 563 samples given at apple_dataset/train
Training on: ['apple', 'damaged_apple']
Training with Batch Size: 8
Number of Training Samples: 563
Number of Validation Samples: 150
Number of Experiments: 50
Training with transfer learning from pretrained Model
WARNING:tensorflow:period argument is deprecated. Please use 'save_freq' to specify the frequency in number of batches seen.
WARNING:tensorflow:os.environ argument is deprecated and will be removed, use 'min_delta' instead.
/usr/local/lib/python3.6/dist-packages/tensorflow/python/keras/engine/training.py:1844: UserWarning: 'Model.fit_generator' is deprecated and will be removed in a future version. Please use 'Model.fit' instead.
warnings.warn("'Model.fit_generator' is deprecated and will be removed in a future version. Please use 'Model.fit' instead.")
WARNING:tensorflow:Model failed to serialize as JSON. Ignoring... Layer YoloLayer has arguments in '__init__' and therefore must override 'get_config'.
/usr/local/lib/python3.6/dist-packages/tensorflow/python/data/ops/dataset_ops.py:3504: UserWarning: Even though the tf.config.experimental_run_functions_eagerly option is set, this operation will not be run eagerly.
warnings.warn("Even though the tf.config.experimental_run_functions_eagerly option is set, this operation will not be run eagerly.")
Epoch 1/50
568/568 [=====] - 484s 831ms/step - loss: 96.4445 - yolo_layer_6_loss: 12.4051 - yolo_layer_7_loss: 27.1762 - yolo_layer_8_loss: 45.3178 - val_loss: 53.8468
Epoch 2/50
568/568 [=====] - 429s 754ms/step - loss: 48.3179 - yolo_layer_6_loss: 6.1802 - yolo_layer_7_loss: 11.8629 - yolo_layer_8_loss: 19.3310 - val_loss: 46.9162
Epoch 3/50
568/568 [=====] - 437s 768ms/step - loss: 41.3287 - yolo_layer_6_loss: 4.7915 - yolo_layer_7_loss: 9.7666 - yolo_layer_8_loss: 16.5045 - val_loss: 38.9622
Epoch 4/50
568/568 [=====] - 432s 760ms/step - loss: 37.9825 - yolo_layer_6_loss: 4.4229 - yolo_layer_7_loss: 8.8794 - yolo_layer_8_loss: 14.9790 - val_loss: 35.8438
Epoch 5/50
568/568 [=====] - 438s 771ms/step - loss: 35.4611 - yolo_layer_6_loss: 4.0843 - yolo_layer_7_loss: 8.0527 - yolo_layer_8_loss: 14.1057 - val_loss: 35.4190
Epoch 6/50
568/568 [=====] - 440s 780ms/step - loss: 34.7106 - yolo_layer_6_loss: 4.4701 - yolo_layer_7_loss: 8.0361 - yolo_layer_8_loss: 13.4007 - val_loss: 34.5009

```

0.3.3 Step 3 — Predicting using our trained Model



As we can see, the trained apple detection model is able to detect all the apples in the image as well as identify the apple with a defect (apple with a spot). You can use this trained model to count apples

```
count=0;
for i in detections:
    count=count+1;
for detection in detections:
    print(detection["name"], " : ", detection["percentage_probability"], " : ", detection["box_points"],)

print("the total number of apples = ", count);

apple : 88.6323094367981 : [31, 10, 196, 137]
apple : 88.17557096481323 : [204, 5, 379, 133]
apple : 81.28604292869568 : [224, 217, 402, 354]
apple : 81.59140944480896 : [31, 247, 225, 356]
apple : 71.5072512626648 : [0, 0, 90, 40]
apple : 60.85934638977051 : [116, 0, 250, 38]
damaged_apple : 75.5850613117218 : [132, 109, 271, 250]
apple : 85.36276817321777 : [300, 105, 411, 254]
apple : 85.44551730155945 : [0, 127, 122, 281]
apple : 84.3239963054657 : [386, 0, 412, 113]
apple : 85.06371974945068 : [0, 11, 24, 129]
the total number of apples = 11
```

The complete code along with other predicted images can be accessed here:

https://github.com/Tajamul21/Detection-Classification-and-Semantic_Segmentation-of-apples

0.4 Semantic Segmentation

The goal of semantic image segmentation is to label each pixel of an image with a corresponding class of what is being represented. Because we're predicting for every pixel in the image, this task is commonly referred to as dense prediction.

Unlike the previous tasks, the expected output in semantic segmentation are not just labels and bounding box parameters. The output itself is a high resolution image (typically of the same size as input image) in which each pixel is classified to a particular class. Thus it is a pixel level image classification.

The different operations that are typically used in a Convolutional Network. Please make a note of the terminologies used.

i. Convolution operation

There are two inputs to a convolutional operation i) A 3D volume (input image) of size ($n_{in} \times n_{in} \times \text{channels}$)

ii) A set of 'k' filters (also called as kernels or feature extractors) each one of size ($f \times f \times \text{channels}$), where f is typically 3 or 5.

The output of a convolutional operation is also a 3D volume (also called as output image or feature map) of size ($n_{out} \times n_{out} \times k$).

The relationship between n_{in} and n_{out} is as follows:

$$n_{out} = \left\lfloor \frac{n_{in} + 2p - k}{s} \right\rfloor + 1$$

n_{in} : number of input features


n_{out} : number of output features

k : convolution kernel size

p : convolution padding size

s : convolution stride size

The function of pooling is to reduce the size of the feature map so that we have fewer parameters in the network. the size of the filter and strides are two important hyper-parameters in the max pooling operation.



0.4.1 U-Net

U-Net is an architecture for semantic segmentation. It consists of a contracting path and an expansive path. The contracting path follows the typical architecture of a convolutional network. It consists of the repeated application of two 3x3 convolutions (unpadded convolutions), each followed by a rectified linear unit (ReLU) and a 2x2 max pooling operation with stride 2 for downsampling. At each downsampling step we double the number of feature channels. Every step in the expansive path consists of an upsampling of the feature map followed by a 2x2 convolution (“up-convolution”) that halves the number of feature channels, a concatenation with the correspondingly cropped feature map from the contracting path, and two 3x3 convolutions, each followed by a ReLU. The cropping is necessary due to the loss of border pixels in every convolution. At the final layer a 1x1 convolution is used to map each 64-component feature vector to the desired number of classes. In total the network has 23 convolutional layers.

```
def get_unet(input_img, n_filters = 16, dropout = 0.1, batchnorm = True):
    """Function to define the UNET Model"""
    # Contracting Path
    c1 = conv2d_block(input_img, n_filters * 1, kernel_size = 3, batchnorm = batchnorm)
    p1 = MaxPooling2D((2, 2))(c1)
    p1 = Dropout(dropout)(p1)

    c2 = conv2d_block(p1, n_filters * 2, kernel_size = 3, batchnorm = batchnorm)
    p2 = MaxPooling2D((2, 2))(c2)
    p2 = Dropout(dropout)(p2)

    c3 = conv2d_block(p2, n_filters * 4, kernel_size = 3, batchnorm = batchnorm)
    p3 = MaxPooling2D((2, 2))(c3)
    p3 = Dropout(dropout)(p3)

    c4 = conv2d_block(p3, n_filters * 8, kernel_size = 3, batchnorm = batchnorm)
    p4 = MaxPooling2D((2, 2))(c4)
    p4 = Dropout(dropout)(p4)

    c5 = conv2d_block(p4, n_filters * 16, kernel_size = 3, batchnorm = batchnorm)

    # Expansive Path
    u6 = Conv2DTranspose(n_filters * 8, (3, 3), strides = (2, 2), padding = 'same')(c5)
    u6 = concatenate([u6, c4])
    u6 = Dropout(dropout)(u6)
    c6 = conv2d_block(u6, n_filters * 8, kernel_size = 3, batchnorm = batchnorm)

    u7 = Conv2DTranspose(n_filters * 4, (3, 3), strides = (2, 2), padding = 'same')(c6)
    u7 = concatenate([u7, c3])
    u7 = Dropout(dropout)(u7)
    c7 = conv2d_block(u7, n_filters * 4, kernel_size = 3, batchnorm = batchnorm)

    u8 = Conv2DTranspose(n_filters * 2, (3, 3), strides = (2, 2), padding = 'same')(c7)
    u8 = concatenate([u8, c2])
    u8 = Dropout(dropout)(u8)
    c8 = conv2d_block(u8, n_filters * 2, kernel_size = 3, batchnorm = batchnorm)

    u9 = Conv2DTranspose(n_filters * 1, (3, 3), strides = (2, 2), padding = 'same')(c8)
    u9 = concatenate([u9, c1])
    u9 = Dropout(dropout)(u9)
    c9 = conv2d_block(u9, n_filters * 1, kernel_size = 3, batchnorm = batchnorm)

    outputs = Conv2D(1, (1, 1), activation='sigmoid')(c9)
    model = Model(inputs=[input_img], outputs=[outputs])
    return model
```

0.4.2 Training

Model is compiled with Adam optimizer and we use binary cross entropy loss function since there are only two classes (apple and no apple). We use Keras callbacks to implement:

Learning rate decay if the validation loss does not improve for 5 continues epochs.

Early stopping if the validation loss does not improve for 10 continues epochs.

Save the weights only if there is improvement in validation loss.

We use a batch size of 32.

Help Save failed

+ Code + Text

RAM Disk

Editing

[14] model.summary()

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
img (InputLayer)	[(None, 128, 128, 1) 0		
conv2d_1 (Conv2D)	(None, 128, 128, 16) 160		img[0][0]
batch_normalization_1 (BatchNor	(None, 128, 128, 16) 64		conv2d_1[0][0]
activation_1 (Activation)	(None, 128, 128, 16) 0		batch_normalization_1[0][0]
max_pooling2d (MaxPooling2D)	(None, 64, 64, 16) 0		activation_1[0][0]
dropout (Dropout)	(None, 64, 64, 16) 0		max_pooling2d[0][0]
conv2d_3 (Conv2D)	(None, 64, 64, 32) 4640		dropout[0][0]
batch_normalization_3 (BatchNor	(None, 64, 64, 32) 128		conv2d_3[0][0]
activation_3 (Activation)	(None, 64, 64, 32) 0		batch_normalization_3[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 32) 0		activation_3[0][0]
dropout_1 (Dropout)	(None, 32, 32, 32) 0		max_pooling2d_1[0][0]
conv2d_5 (Conv2D)	(None, 32, 32, 64) 18496		dropout_1[0][0]
batch_normalization_5 (BatchNor	(None, 32, 32, 64) 256		conv2d_5[0][0]
activation_5 (Activation)	(None, 32, 32, 64) 0		batch_normalization_5[0][0]
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 64) 0		activation_5[0][0]
dropout_2 (Dropout)	(None, 16, 16, 64) 0		max_pooling2d_2[0][0]
conv2d_7 (Conv2D)	(None, 16, 16, 128) 73856		dropout_2[0][0]
batch_normalization_7 (BatchNor	(None, 16, 16, 128) 512		conv2d_7[0][0]

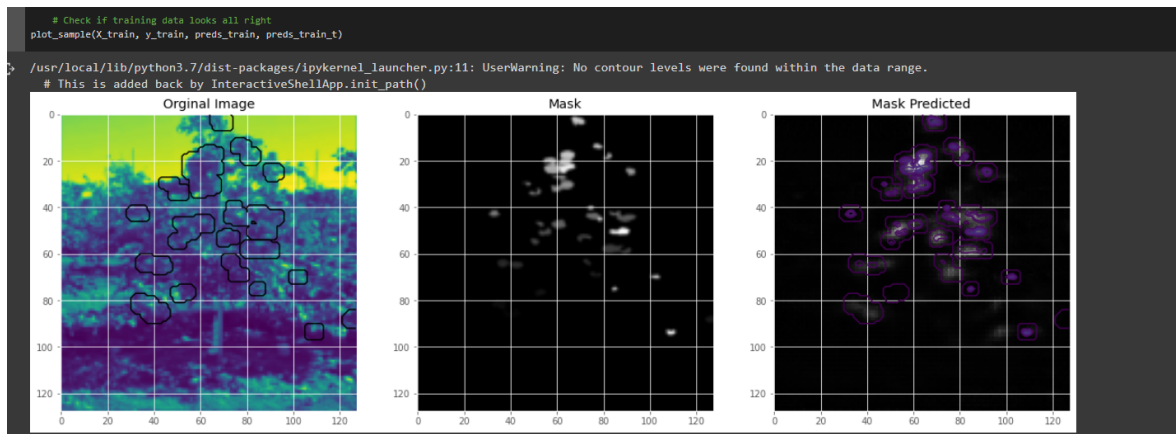
0.4.3 Inference

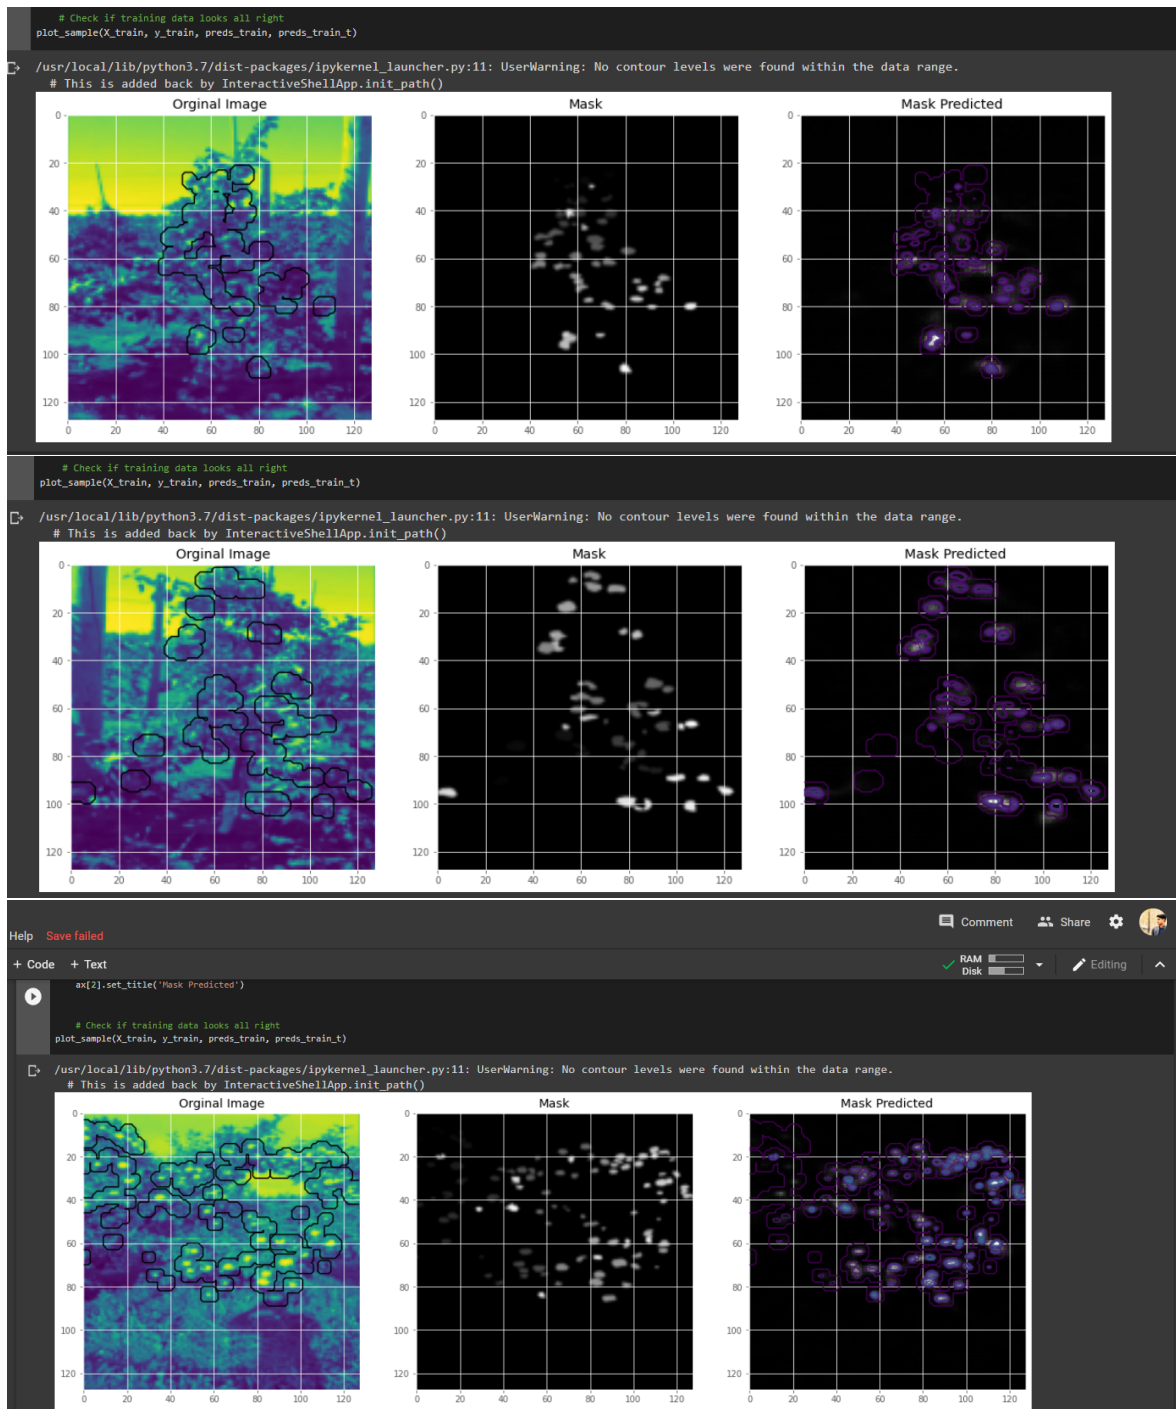
For each pixel we get a value between 0 to 1.

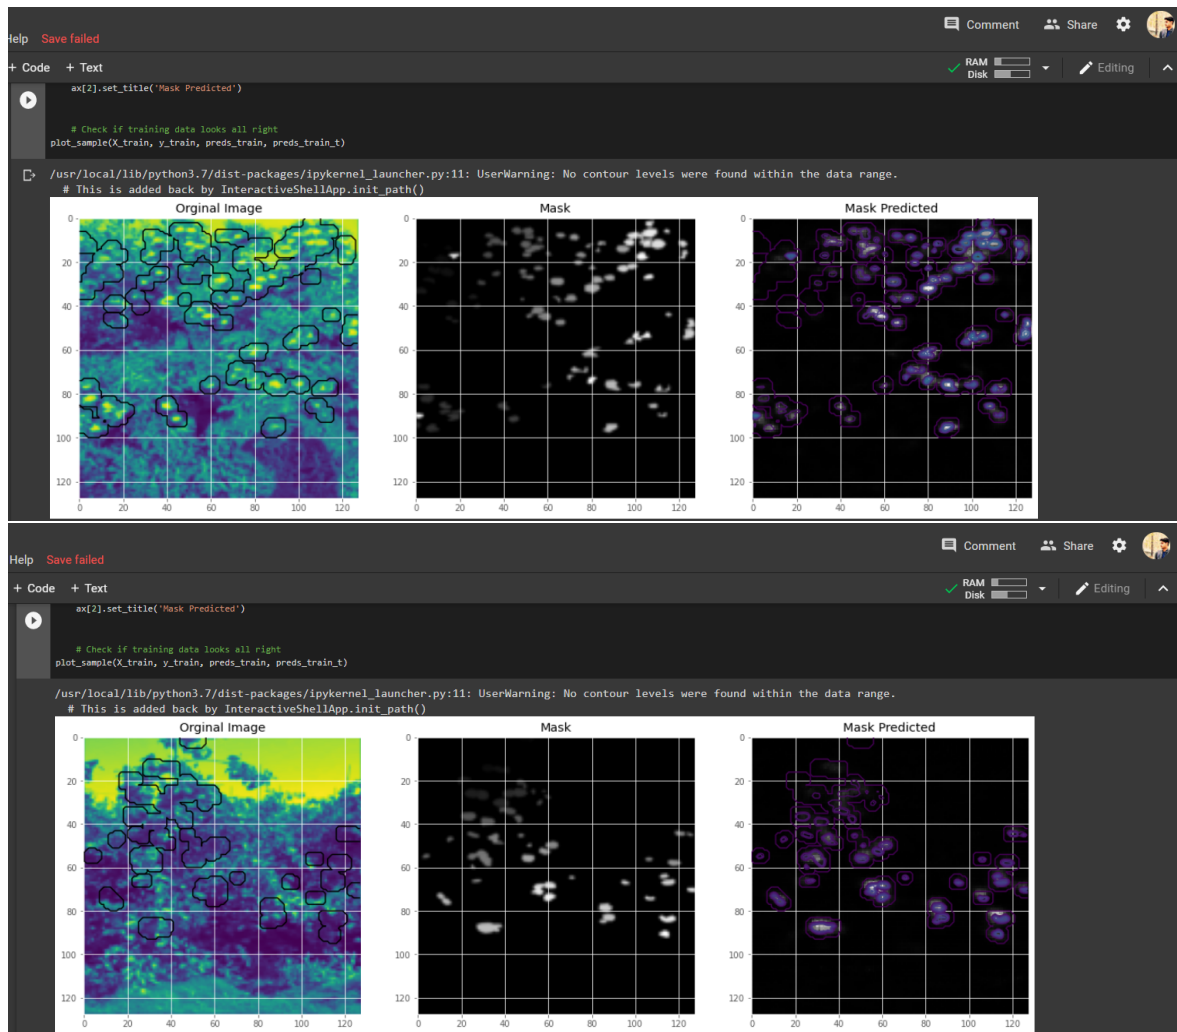
0 represents no apple and 1 represents apple.

We take 0.5 as the threshold to decide whether to classify a pixel as 0 or 1.

0.4.4 Our Results on training set







The complete code along with other predicted(segmented) images can be accessed here:

<https://github.com/Tajamul21/Detection-Classification-and-Semantic-Segmentation-of-apples>